

Get started

Open in app



490K Followers · About Follow

You have **1** free member-only story left this month. [Sign up for Medium and get an extra one](#)



Photo by [AbsolutVision](#) on [Unsplash](#)

Create A Simple Search Engine Using Python

Information retrieval using cosine similarity and term-document matrix with TF-IDF weighting.



Irfan Alghani Khalid Sep 21 · 7 min read ★

All of us have used a search engine, in example Google, in every single day for searching everything, even on simple things. But have you ever imagined, how that search engine can retrieve all of our documents based on what we want to search (query)?

In this article, I will show you on how to build a simple search engine from scratch using Python and its supporting library. After you read the article, I hope you can understand how to build your own search engine based on what you need. Without further, let's go!

Side note: I've also created a notebook of the code, so if you want to follow along with me you can click on this link [here](#). Also, the documents that I will use is in Indonesian. But don't worry, you can use any documents regardless of the language.

Outline

Before we get our hands dirty, let me give you the steps on how to implement this, and on each section, I will explain on how to build it. They are,

- Preparing the documents
- Create a Term-Document Matrix with TF-IDF weighting
- Calculate the similarities between query and documents using Cosine Similarity
- Retrieve the articles that have the highest similarity on it.

The Process

Retrieve the documents

The first thing that we have to do is to retrieve the documents from the Internet. In this case, we can use web scraping to extract documents from a website. I will scrape documents from kompas.com on sport category, especially on the popular articles. Because of the documents are using HTML format, we initialize a BeautifulSoup object to parse the HTML file, so we can extract each element that we want much easier.

Based on the figures below, I've shown the screenshot of the website with an inspect element to it. On Figure 1, I've shown the tags that we want to retrieve, which is the href attribute of the highlighted <a> tag with class "most__link". On the Figure 2, We will retrieve text on <p> tags from <div> tag with class "read__content".

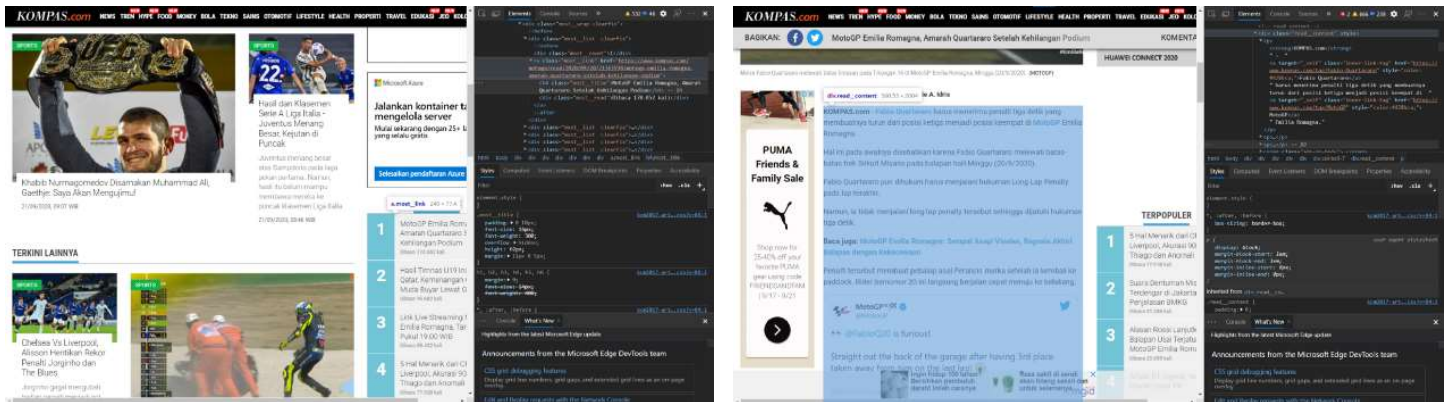


Figure 1, Figure 2

Here is the code that I used for extracting the documents and its explanations on each line,

```
import requests
from bs4 import BeautifulSoup

# Make a request to the website
r = requests.get('https://bola.kompas.com/')

# Create an object to parse the HTML format
soup = BeautifulSoup(r.content, 'html.parser')

# Retrieve all popular news links (Fig. 1)
link = []
for i in soup.find('div', {'class': 'most__wrap'}).find_all('a'):
    i['href'] = i['href'] + '?page=all'
    link.append(i['href'])

# For each link, we retrieve paragraphs from it, combine each
paragraph as one string, and save it to documents (Fig. 2)
documents = []
for i in link:
    # Make a request to the link
    r = requests.get(i)

    # Initialize BeautifulSoup object to parse the content
    soup = BeautifulSoup(r.content, 'html.parser')

    # Retrieve all paragraphs and combine it as one
    sen = []
    for i in soup.find('div',
```

```
{'class':'read__content'}).find_all('p'):
    sen.append(i.text)

# Add the combined paragraphs to documents
documents.append(' '.join(sen))
```

Clean the documents

Right after we extract the documents, we have to clean it, so our retrieval process becomes much easier. For each document, we have to remove all unnecessary words, numbers and punctuations, lowercase the word, and remove the doubled space. Here is the code for it,

```
import re

documents_clean = []
for d in documents:
    # Remove Unicode
    document_test = re.sub(r'^\x00-\x7F+', ' ', d)
    # Remove Mentions
    document_test = re.sub(r'@\w+', '', document_test)
    # Lowercase the document
    document_test = document_test.lower()
    # Remove punctuations
    document_test = re.sub(r'[%s]' % re.escape(string.punctuation),
    ' ', document_test)
    # Lowercase the numbers
    document_test = re.sub(r'[0-9]', '', document_test)
    # Remove the doubled space
    document_test = re.sub(r'\s{2,}', ' ', document_test)
    documents_clean.append(document_test)
```

Create Term-Document Matrix with TF-IDF weighting

After each document is clean, it's time to create a matrix. Thankfully, scikit-learn library has prepared for us the code of it, so we don't have to implement it from scratch. The code looks like this,

```
from sklearn.feature_extraction.text import TfidfVectorizer

# Instantiate a TfidfVectorizer object
vectorizer = TfidfVectorizer()

# It fits the data and transform it as a vector
X = vectorizer.fit_transform(docs)
```

```
# Convert the X as transposed matrix
X = X.T.toarray()

# Create a DataFrame and set the vocabulary as the index
df = pd.DataFrame(X, index=vectorizer.get_feature_names())
```

The result (matrix) will become a representation of the documents. By using that, we can find the similarity between different documents based on the matrix. The matrix looks like this,

	0	1	2	3	4	5	6	7	8	9
abdulkarim	0.000000	0.022825	0.000000	0.0	0.000000	0.000000	0.0	0.000000	0.000000	0.000000
abraham	0.000000	0.000000	0.000000	0.0	0.000000	0.000000	0.0	0.000000	0.060882	0.000000
absen	0.000000	0.000000	0.000000	0.0	0.079834	0.000000	0.0	0.000000	0.000000	0.000000
ada	0.075228	0.013554	0.040336	0.0	0.000000	0.052744	0.0	0.029377	0.000000	0.000000
adalah	0.022669	0.000000	0.060775	0.0	0.064286	0.047682	0.0	0.026558	0.000000	0.031469

Term-Document Matrix

The matrix above is called as Term-Document Matrix. It consists of rows that represent by each token (term) from all documents, and the columns consist of the identifier of the document. Inside of the cell is the number of frequency of each word that is weighted by some number.

We will use the column vector, which is a vector that represents each document to calculate the similarity with a given query. We can call this vector as **embeddings**.

For calculating the cell value, the code uses the TF-IDF method to do this. TF-IDF (Term Frequency — Inverse Document Frequency) is a frequency of a word that is weighted by IDF. Let me explain each one of them,

Term Frequency (TF) is a frequency of term (t) on document (d). The formula looks like this,

$$tf_{t,d} = count(t, d)$$

Beside of that, we can use a log with bases of 10 to calculate the TF, so the number becomes smaller, and the computation process becomes faster. Also, make sure to add

one on it because we don't want log 0 exist.

$$tf_{t,d} = \log_{10}(\text{count}(t, d) + 1)$$

Then, there is the **Inverse Document Frequency (IDF)**. This formula will be used for calculating the rarity of the word in all documents. It will be used as weights for the TF. If a word is frequent, then the IDF will be smaller. In opposite, if the word is less frequent, then the IDF will be larger. The formula looks like this,

$$idf_t = \log_{10}\left(\frac{N}{df_t}\right)$$

Recall the TF-IDF, we can see how does it affect the value on each cell. It will remove all the words that are frequently shown in documents but at the same time not important, such as and, or, even, actually, etc. Based on that, we use this as the value on each cell on our matrix.

Calculate the similarity using cosine similarity.

After we create the matrix, we can prepare our query to find articles based on the highest similarity between the document and the query. To calculate the similarity, we can use the cosine similarity formula to do this. It looks like this,

$$\text{cosine}(q, d) = \frac{q * d}{|q||d|}$$

The formula calculates the dot product divided by the multiplication of the length on each vector. The value ranges from [1, 0], but in general, the cosine value ranges from [-1, 1]. Because there are no negative values on it, we can ignore the negative value because it never happens.

Now, we will implement the code to find similarities on documents based on a query. The first thing that we have to do is to transform the query as a vector on the matrix that we have. Then, we calculate the similarities between them. And finally, we retrieve all documents that have values above 0 in similarity. The code looks like this,

```

def get_similar_articles(q, df):
    print("query:", q)
    print("Berikut artikel dengan nilai cosine similarity tertinggi:
")

    # Convert the query become a vector
    q = [q]
    q_vec = vectorizer.transform(q).toarray().reshape(df.shape[0],)
    sim = {}

    # Calculate the similarity
    for i in range(10):
        sim[i] = np.dot(df.loc[:, i].values, q_vec) /
np.linalg.norm(df.loc[:, i]) * np.linalg.norm(q_vec)

    # Sort the values
    sim_sorted = sorted(sim.items(), key=lambda x: x[1], reverse=True)

    # Print the articles and their similarity values
    for k, v in sim_sorted:
        if v != 0.0:
            print("Nilai Similaritas:", v)
            print(docs[k])
            print()

# Add The Query
q1 = 'barcelona'

# Call the function
get_similar_articles(q1, df)

```

Suppose that we want to find articles that talk about Barcelona. If we run the code based on that, we will get the result like this,

```

query: barcelona
Berikut artikel dengan nilai cosine similarity tertinggi:
Nilai Similaritas: 0.4641990113096689
    kompas.com perombakan skuad yang dilakukan pelatih anyar barcelona
ronald koeman memakan korban baru terkini ronald koeman dikabarkan
akan mendepak bintang muda barcelona yang baru berusia tahun riqui
puig menurut media spanyol rac koeman sudah meminta riqui puig
mencari tim baru karena tidak masuk dalam rencananya di barcelona
rumor itu semakin kuat karena puig....

Nilai Similaritas: 0.4254860197361395
    kompas.com pertandingan trofeo joan gamper mempertemukan barcelona
dengan salah satu tim promosi liga spanyol elche laga barcelona vs
elche usai digelar di camp nou pada minggu dini hari wib trofeo joan
gamper merupakan laga tahunan yang diadakan oleh barca kali ini
sudah memasuki edisi ke blaugrana julukan tuan rumah menang dengan
skor gol kemenangan barcelona....

```

Final Thoughts

That is how we can create a simple search engine using Python and its dependencies. It still very basic, but I hope you can learn something from here and can implement your own search engine based on what you need. Thank you.

References

[1] Jurafsky, D. & Martin, J.H. *Speech and Language Processing* (2000), Prentice Hall.

Sign up for The Daily Pick

By Towards Data Science

Hands-on real-world examples, research, tutorials, and cutting-edge techniques delivered Monday to Thursday. Make learning your daily ritual. [Take a look](#)

Your email

Get this newsletter

By signing up, you will create a Medium account if you don't already have one. Review our [Privacy Policy](#) for more information about our privacy practices.

Data Science

Information Retrieval

NLP

Programming

Python

[About](#) [Help](#) [Legal](#)

Get the Medium app



